

Alex Rodriguez

Email: alex.rodriguez@techmail.com

Phone: +1 (212) 555-0199

Location: New York, NY

LinkedIn: linkedin.com/in/alex-rodriguez-data

PROFESSIONAL SUMMARY

Senior Data Engineer with 7+ years of experience specializing in building high-throughput, real-time data pipelines. Expertise in distributed systems like Kafka and Spark, and modern data warehousing using Snowflake. Proven ability to deliver scalable infrastructure on AWS.

EDUCATION

Master of Science in Computer Engineering

New York University, 2017

Bachelor of Science in Information Technology

Rutgers University, 2015

CERTIFICATIONS

- AWS Certified Data Analytics - Specialty
- Snowflake Certified Advanced Data Engineer

WORK EXPERIENCE

Lead Data Engineer | JPMorgan Chase | February 2020 - Present

- Architected and implemented a real-time fraud detection pipeline using Apache Kafka and Spark Streaming, processing 5M events/second.
- Designed and migrated legacy ETL jobs to a fully serverless, Python-based Airflow environment on AWS, reducing operational costs by 30%.

- Managed and optimized a 50+ TB Snowflake data warehouse, improving critical query performance by 45% through clustering and materialized views.
- Mentored a team of 4 junior engineers on data modeling (Dimensional Modeling) and pipeline best practices.

Data Engineer | Bloomberg | August 2017 - January 2020

- Developed and maintained batch processing workflows using Python, PySpark, and AWS Glue for market data ingestion.
- Implemented data governance policies, achieving 99.9% data quality for financial reporting systems.
- Managed a Hadoop ecosystem, including Hive and HDFS, for large-scale data lake operations.

TECHNICAL SKILLS

Programming Languages: Python (Expert), SQL (Expert), Scala (Intermediate)

Big Data & Streaming: Apache Spark, Kafka, Hadoop, Hive

Data Warehousing: Snowflake, AWS Redshift, PostgreSQL

Cloud Platforms: AWS (Glue, S3, EMR, Lambda)

Orchestration & DevOps: Airflow, Terraform, Docker, Git

PROJECTS

Real-Time Market Data Streamer

- Built a Python-based Kafka consumer application to process stock quotes, feeding a Redshift data warehouse.
- Achieved sub-second latency for critical financial analytics reports.

Data Quality Monitor

- Developed a custom data validation framework in Spark to automatically flag schema and integrity errors.